



Convolutional Neural Networks for Text Classification using Intel Nervana Neon

Kripa Sankaranarayanan, Yinyin Liu

Acknowledgements: Vipul Lal, Hanlin Tang, Sejun Kim

TOPICS

Neon


CNN in Text Classification

Two Implementation examples of CNN for Text classification in Neon



INTEL[®] NERVANA[™] FULL STACK AI PLATFORM

- **Nervana Cloud** → Build an AI POC
- **neon** → Train DL models quickly
- **Intel Nervana Graph** → any framework, any hardware
- **Intel Nervana HW** → industry leading AI, coming soon


“deep learning by design”



neon
deep learning
framework



DEEP LEARNING WITH NEON

 NervanaSystems / neon

Watch 285

★ Unstar 2,357

Fork 504

<> Code

🔔 Issues 22

🔗 Pull requests 2

📁 Projects 0

📖 Wiki

📊 Pulse

📈 Graphs

⚙️ Settings

Fast, scalable, easy-to-use Python based Deep Learning Framework by Nervana™ <http://neon.nervanasys.com/> — Edit

📦 875 commits

🌿 3 branches


📦 22 releases

👤 47 contributors

📄 Apache-2.0

Branch: master ▾ New pull request

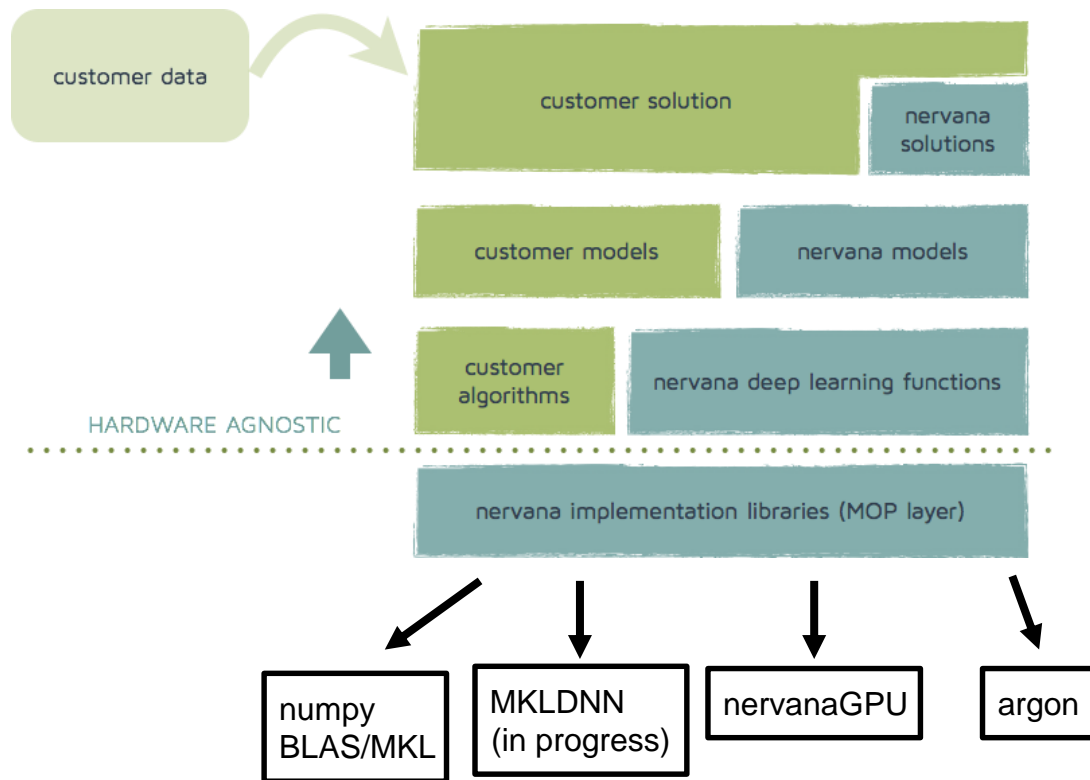
Create new file Upload files Find file Clone or download ▾

 jennifermeyers Update for release 1.6.0 Latest commit 4fb5ff6 23 days ago


bin	Updates for Kepler and Pascal GPU support; Examples doc strings updates	4 months ago
doc	Update for release 1.6.0	22 days ago
examples	bugfix and interface cleanup for seq2seq	22 days ago
loader	Fix compiler warnings	22 days ago
neon	bugfix and interface cleanup for seq2seq	22 days ago
tests	NMS kernel with unit tests, proposal layer cleanup, faster rcnn code ...	a month ago

<https://github.com/NervanaSystems/neon>

MULTIPLE COMPUTE BACKENDS



SPEEDING UP DEEP LEARNING



<https://www.nervanasys.com/winograd/>

CURATED MODELS

- <https://github.com/NervanaSystems/ModelZoo>
- Pre-trained weights and models

Single Shot Detection Deep Residual Net

SegNet **Video Activity Detection**


Skip-thought imdb Sentiment Analysis

Deep Speech 2 **LSTM Image Captioning**


Autoencoders Deep Reinforcement Learning

Deep Dream **Fast-RCNN Object Localization**

Generative Adversarial Networks



NERVANA MODEL LIBRARY



Object localization

Fast-RCNN




Image classification

Deep Residual Net



Sentiment analysis

LSTM



Question answering

GRU



Video activity recognition

3D Convolutional Net



Scene classification

Deep Residual Net




Speech recognition

Deep Speech 2

EXAMPLE CURATED MODEL



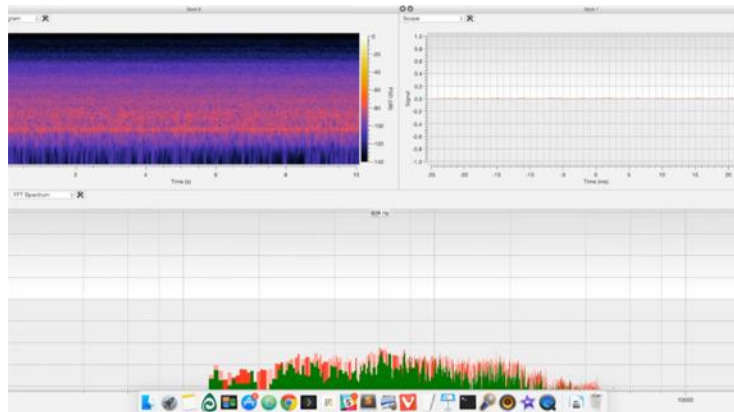
SegNet Model Prediction



- Image segmentation model
- Neon model 5.5x faster than Caffe counterpart
- ~50 fps real time analysis

	Neon (ms)	Caffe (ms)	Speed-up
Forward	101	719	7.1x
Backward	164	746	4.5x
Total	265	1455	5.5x

SPEECH RECOGNITION



Reference	CER	WER	WER (LM)
Hannun-Maas (2014)	10.7	35.8	14.1
Graves-Jaitly (ICML 2014)	9.2	30.1	10.4
Hwang-Sung (ICML 2016)	10.6	38.4	8.8
Miao et al (Eesen)	N/A	N/A	9.1
Bahdanau et al (2016)	6.4	18.6	10.8
Nervana Speech	8.6	32.5	8.4
Baidu DS2 (trained on 12,000 hours)			3.6

Wall Street Journal evaluation data (WSJ0)

NEON WORKFLOW

Generate backend

Load data

Specify model architecture

Define training parameters

Train model

Evaluate

NEON OVERVIEW

Backend	NervanaGPU, NervanaCPU, NervanaMGPU
Datasets	MNIST, CIFAR-10, Imagenet 1K, PASCAL VOC, Mini-Places2, IMDB, Penn Treebank, Shakespeare Text, bAbI, Hutter-prize, UCF101, flickr8k, flickr30k, COCO
Initializers	Constant, Uniform, Gaussian, Glorot Uniform, Xavier, Kaiming, IdentityInit, Orthonormal
Optimizers	Gradient Descent with Momentum, RMSProp, AdaDelta, Adam, Adagrad, MultiOptimizer
Activations	Rectified Linear, Softmax, Tanh, Logistic, Identity, ExpLin
Layers	Linear, Convolution, Pooling, Deconvolution, Dropout, Recurrent, Long Short-Term Memory, Gated Recurrent Unit, BatchNorm, LookupTable, Local Response Normalization, Bidirectional-RNN, Bidirectional-LSTM
Costs	Binary Cross Entropy, Multiclass Cross Entropy, Sum of Squares Error
Metrics	Misclassification (Top1, TopK), LogLoss, Accuracy, PrecisionRecall, ObjectDetection

CNN IN TEXT CLASSIFICATION

Labels to text – sentiment, topic

Classifiers

- N-grams with linear classifier
- Deep learning : RNN based on LSTM, GRU
- CNN
 - Reduction in computation, parallelization
 - Hierarchical vs Sequential for RNN
 - Better with large scale datasets
 - No semantic or syntactic knowledge of language

CNN IN TEXT CLASSIFICATION

Character level classification

- Text treated as raw signal at character level
- Characters are a necessary construct - work for different languages
- Characters – one hot representation

Sentence level classification

- Word vectors

EXAMPLE 1:

CHARACTER LEVEL TEXT CLASSIFICATION WITH CNN

DATA

That's right....the red velvet cake.....ohh
this stuff is so good!



POSITIVE

Then, as if I hadn't wasted enough of my life there, they
poured salt in the wound by drawing out the time it took to
bring the check



NEGATIVE

INPUTS

- Sequence of encoded characters
- Encoding
 - prescribe an alphabet of size m for the input language
 - quantize each character using 1-of- m encoding (or “one-hot” encoding)
 - sequence of characters is transformed to a sequence of m sized vectors with fixed length l_0 .

DATA PREPROCESSING

“That's right....the red velvet cake.....ohh
this stuff is so good!”



[50, 38, 30, 50, 7, 47, 3, 48, 37,
35, 38, 50, 16, 16, 16, 16, 50,
38, 33, 3, 48, 33, 34, 3, 52, 33,
42, 52, 33, 50, 3, 31, 30, 39, 33,
16, 16, 16, 16, 16, 43, 38, 38,
38, 3, 50, 38, 37, 47, 3, 47, 50,
49, 36, 36, 3, 37, 47, 3, 47, 43,
3, 35, 43, 43, 34, 16]

- Number of characters limited to 56 (vocab size)
- Truncate each sample to 256 characters [from the left]
- data format is:
sentences * max characters per sentence * vocab size
(#Samples * 256 * 56)


MODEL

- Layers:
 - Convolutional module, computes a 1-D convolution
 - Max-pooling
 - Fully connected
- Algorithm: stochastic gradient descent (SGD)
- Activation: ReLU
- Initialization : Gaussian

MODEL

Input
[50, 38, 30,
50, 7, 47,
3, 48, 37,
35, 38, 50,
16, 16, 16,
16, 50, ...,
43, 43, 34,
16]

Character
Quantization




256 X 56

Convolution and Pooling
Layers

$$\text{Convolution: } h(y) = \sum_{x=0}^k f(x) \cdot g(y \cdot d - x + c)$$

$$\text{Pooling: } h(y) = \max g(y \cdot d - x + c)$$



Fully Connected
Layers

Class 0

Class 1

NEON COMMANDS

Object that returns minibatches of data (nsamples, 256, 56) (nsamples, 2)

```
train_set = ArrayIterator(X_train, y_train, nclass=nclass)
valid_set = ArrayIterator(X_test, y_test, nclass=nclass)

init = Gaussian(loc=0.0, scale=0.05)

layers = [ Conv((5, 5, f_size), init=init, activation=relu), Pooling(3, op='max'),
           Conv((5, 5, f_size), init=init, activation=relu), Pooling(3, op='max'),
           Conv((3, 3, f_size), init=init, activation=relu),
           Conv((3, 3, f_size), init=init, activation=relu),
           Dropout(keep=.5),
           Linear(nout=lin_f, init=init)]

char_cnn = Model(layers=layers)

cost = GeneralizedCost(costfunc=CrossEntropyBinary())
optimizer = GradientDescentMomentum(0.01, momentum_coef=0.9)

callbacks = Callbacks(char_cnn, eval_set=valid_set)

char_cnn.fit(train_set, optimizer=optimizer, num_epochs=args.epochs, cost=cost,
             callbacks=callbacks)
```

Set up initializer

Set up Layers

EXAMPLE 2:

SENTENCE CLASSIFICATION WITH CNN

DATASETS

Sequence of encoded words

Sentence max words 56

Embedding size 300

Vocab size ~20K

nsamples: ~9K/1K training/validation

MODEL

- Layers:
 - Embedding Layer
 - 3 Convolutional modules with different filter sizes
 - Max Pooling
 - Dropout layer
 - Affine – fully connected with Softmax
- Algorithm: stochastic gradient descent (SGD)
- Activation: ReLU
- Initialization : Gaussian


MODEL

Data


[751, 1727, 30,
50, 7, 47, 3, 48,
37, 1474, 38,
...36, 36, 17,
137, 47, 47,
8127, 7, 35,
1506, 1834]




Embedding Layer



Convolution and Pooling



Affine



Class 1



Class 0

NEON COMMANDS

Object that returns
minibatches of data

```
train_set = ArrayIterator(X_train, y_train, nclass=nclass)
valid_set = ArrayIterator(X_test, y_test, nclass=nclass)
```

Set up initializer

```
init_uni = Uniform(low=-0.1, high=0.1)
g_uni = GlorotUniform()
```

Set up Layers

```
def gen_layers(fvals):
    (f1, f2, f3) = fvals
    branch1 = [Conv((f1, f1, conv_f), padding = 2, strides = 1, init=init, activation=relu),
    Pooling(3, op='max')]
    branch2 = [Conv((f2, f2, conv_f), padding = 2, strides = 1, init=init, activation=relu),
    Pooling(3, op='max')]
    branch3 = [Conv((f3, f3, conv_f), padding = 2, strides = 1, init=init, activation=relu),
    Pooling(3, op='max')]
    return MergeBroadcast(layers=[branch1, branch2, branch3], merge="stack")

snt_cnn = Model(layers=[
    LookupTable(vocab_size, embedding_dim=embed_size, init=uni),
    gen_layers([6,5,3]),
    Dropout(keep=.5),
    Affine(lin_f, init=init, bias=g_uni, activation=Softmax())])
```

SUMMARY

Neon

CNN for Text classification implementations in Neon

- Performance with large scale datasets

ACKNOWLEDGEMENTS

Intel IT

Intel Nervana Neon team

AIM Conference committee

BACKUP

REFERENCES AND LINKS

www.intelnervana.com

<https://github.com/NervanaSystems/neon>

<https://arxiv.org/pdf/1509.01626.pdf> (Character level CNN for text classification)

<https://arxiv.org/pdf/1408.5882.pdf> (CNN for Sentence classification)

<https://arxiv.org/abs/1702.01923> (Comparative Study of CNN and RNN for NLP)

<https://arxiv.org/pdf/1606.01781.pdf> (Very Deep Convolutional Networks for Text Classification)